

Combining Runtime PM and suspend/resume

Kevin Hilman
khilman@linaro.org

For an Introduction to runtime PM
http://people.linaro.org/~khilman/runtime_PM.html

What ? !!!

Now I have to add
Runtime PM to
my driver?

I thought
suspend/resume
was enough !!





System PM

Runtime PM

Today: Separate worlds

Tomorrow

Happy
together



Runtime PM

- suspend on inactivity
- resume when needed



System PM

- suspend when requested
- resume with system

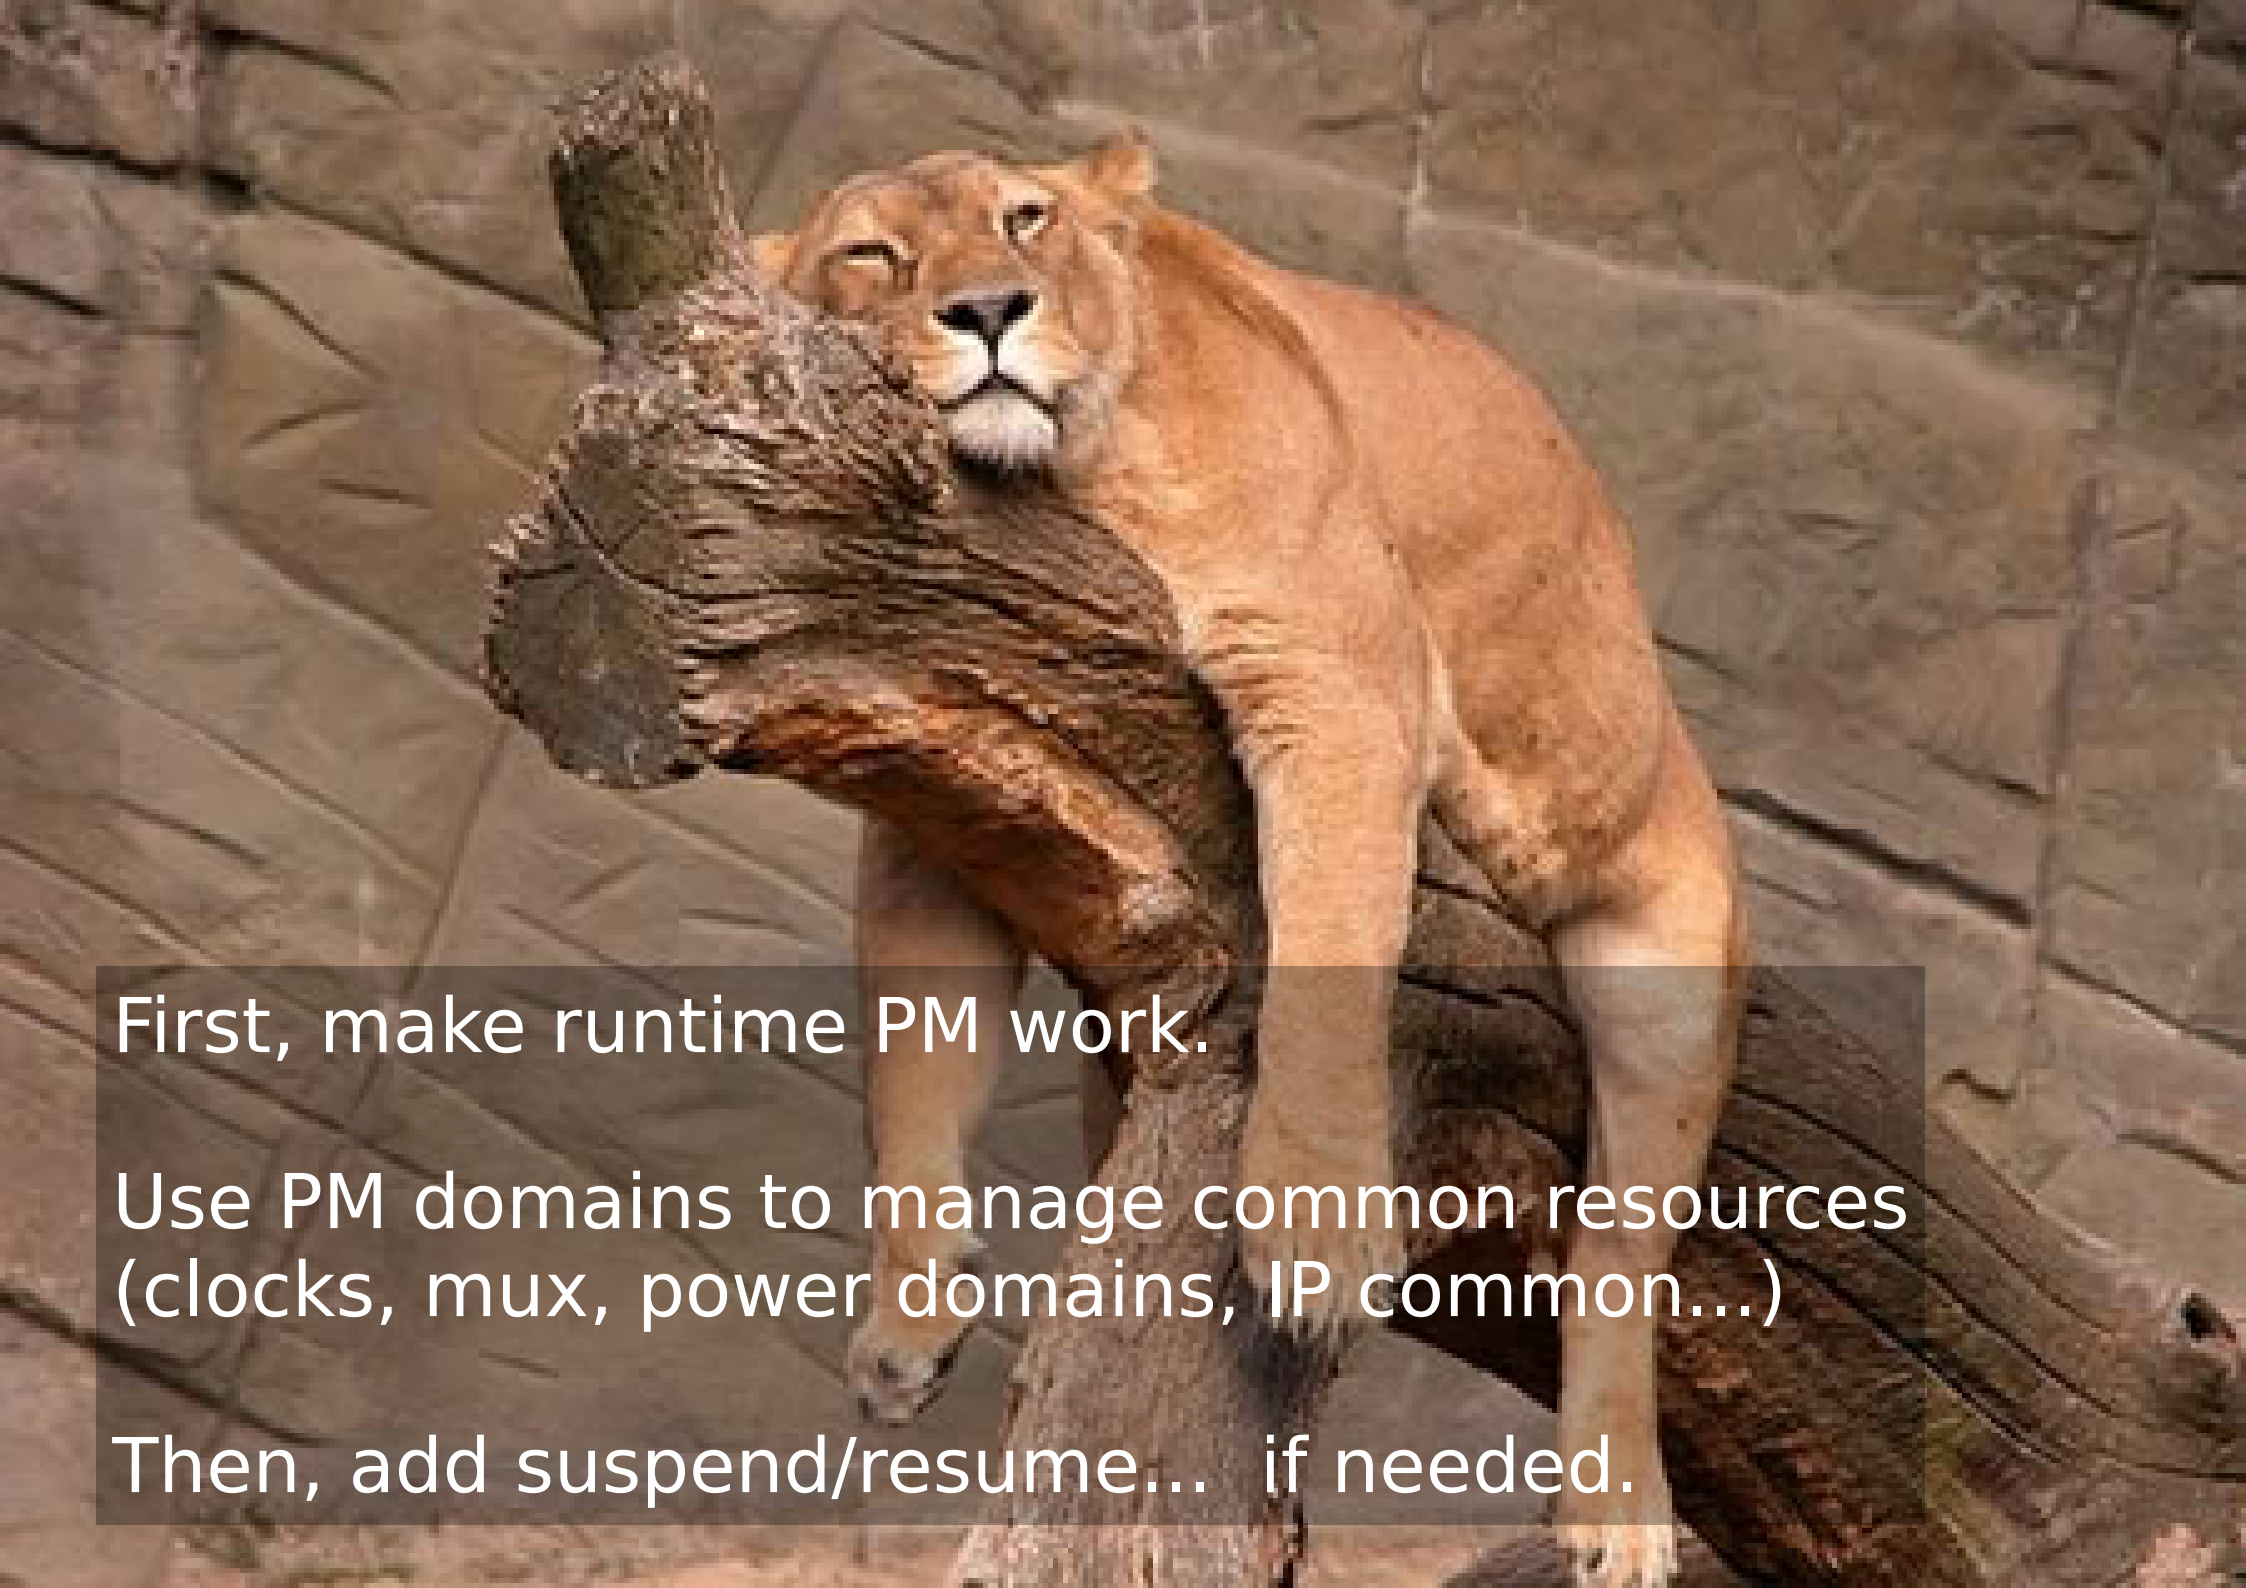


Hey! what if I could
implement runtime PM,
and get suspend
for free?





So... how to combine them?



First, make runtime PM work.

Use PM domains to manage common resources
(clocks, mux, power domains, IP common...)

Then, add suspend/resume... if needed.

Suspend triggers runtime PM

- Force inactivity in suspend
- device is now idle/inactive
- runtime PM takes over



Don't implement suspend

Necessary if device
is used by other
devices during
suspend



Example: OMAP I2C driver

Needed during suspend by other drivers

- I2C-connected RTC
- regulators on I2C-connected PMIC

Solution:

- runtime PM only (no suspend/resume hooks)
- full (re)init per xfer; stateless

Quiz 1: What if I2C driver implemented suspend and was suspended before RTC?

Quiz 2: Who decides suspend ordering?


Great, so...

runtime PM, plus a little "enforcement"
in `->suspend()` gets me both... right?

Yes...

Well...

Almost...

A close-up photograph of a person's green eye. A lit matchstick is held against the upper eyelid, with the red tip of the matchstick touching the skin. The eye is looking directly at the camera. The matchstick is positioned vertically, with the lit tip at the top. The eye is green with a black pupil. The skin around the eye is light-colored. The matchstick is made of wood and has a red tip. The text "Except... there are several ways for devices to be kept awake." is overlaid on the bottom of the image.

Except... there are several
ways for devices to be kept awake.

Tricky: async callbacks

PM workqueue frozen during suspend

- Asynchronous callbacks happen via PM core workqueue (`pm_wq`)
- workqueue frozen during suspend
- Result: no async callbacks

Solution: use `_sync()` methods

Tricky: Autosuspend timeouts

Driver does some work during `->suspend()`

- `_get_sync()`
- `_put_sync_autosuspend()`
- autosuspend callbacks are asynchronous
- Result: device left runtime enabled

Solution: PM domain magic

Tricky: PM core "protections"

PM core: `device_prepare()`

- increments usecount, decrements in `_complete()`
- effectively disables runtime suspend (but not resume)

Example: device runtime suspended: usecount = 0

- PM core prepare: `_get_noresume()`: usecount = 1
- `driver->suspend()`: `_get_sync()`: usecount = 2
- driver uses HW
- `driver->suspend()`: `_put_sync()`: usecount = 1

Result: stuck with usecount > 0, runtime enabled

Solution: PM domain magic

Tricky: userspace disable

runtime PM can be disabled from userspace

```
# echo on > /sys/devices/.../power/control
```

Oops, drivers using only runtime PM will be left runtime enabled during suspend.

Solution: PM domain magic



PM domain magic

Use "late" callbacks

Ensure all devices are suspended

```
struct dev_pm_ops {  
    [...]  
    int (*suspend_late) (struct device *dev);  
    int (*resume_early) (struct device *dev);  
};  
struct dev_pm_domain {  
    struct dev_pm_ops ops;  
};
```

The magic trick

Runtime suspend "enforced" for stuck devices

```
static int my_domain_suspend_late(struct device *dev)>
{
    if (!pm_runtime_status_suspended(dev)) {
        if (pm_generic_runtime_suspend(dev) == 0) {
            /* platform-specific PM */
        }
    }
    return 0;
}
```



Summary

- Runtime PM first
- use PM domains
- Keep drivers generic

What next?

- better understand PM core protections in `device_prepare()`
- Help PM domains proliferate
- What else?

¿ Questions ?

結束